

Become the Maestro of your Genomics Workflow with Bioconductor and Microsoft R Server

Colby Ford

Microsoft R Server is an enterprise-class tool for hosting and managing parallel and distributed workloads of R processes on servers. Organizations that need to process large amounts of data or perform complex processing on the data benefit the most from a parallel architecture like Microsoft R Server. It uses the `RevoScaleR` package, which makes parallelization easy.

In genomics research, we often interact with large amounts of data from complex pipelines in a diverse array of formats. Luckily, [Bioconductor](#) helps make this process simpler by packaging up common sets of processes in ready-to-use R code.

Harnessing the power of both Bioconductor and Microsoft R Server together can help streamline the processing of your genomics data.

All About Bioconductor



Bioconductor is an open-source, open-development software project to provide tools for the analysis and comprehension of high-throughput genomic data. It is based primarily on the R programming language. In other words, it's an extension of R that is specialized for bioinformatics and genomics analyses.

Installation

Once R (either base R or Microsoft R Server) is installed on your local machine, installing Bioconductor is simple. Open R and use the following commands to grab the latest version of Bioconductor.

```
source("https://bioconductor.org/biocLite.R")
biocLite()
```

Once the `biocLite` script has loaded, you can now call any desired packages.

```
library("BiocInstaller")
biocLite("RforProteomics", dependencies = TRUE)
```

Now you are ready to use over 1,000 Bioconductor packages!

Workflows

The field of genomics is very broad, but Bioconductor will often have a solution for every

area. The Bioconductor site is rich with workflow examples to help connect the dots for your research. Check out [Bioconductor's help section](#) for a list of the available workflows.

Here are a few of my favorites:

- [Sequence Analysis](#) - Import fasta, fastq, BAM, gff, bed, wig, and other sequence formats. Trim, transform, align, and manipulate sequences. Perform quality assessment, ChIP-seq, differential expression, RNA-seq, and other workflows.
- [Variant Annotation](#) - Read and write VCF files. Identify structural location of variants and compute amino acid coding changes for non-synonymous variants and predict consequence of amino acid coding changes.
- [High Throughput Assays](#) - Import, transform, edit, analyze and visualize flow cytometric, mass spec, HTqPCR, cell-based, and other assays.
- [Transcription Factor Binding](#) - Find candidate binding sites for known transcription factors via sequence matching.
- [Cancer Genomics](#) - Download, process, and prepare [TCGA](#), [ENCODE](#), and [Roadmap](#) data to interrogate the epigenome of cultured cancer cell lines as well as normal and tumor tissues with high genomic resolution.

Package Database

Didn't see a workflow that exactly fit your needs? No problem! Bioconductor has a nice list of available packages that will help you find the right one for the problem at hand. Using the search box shown below, I searched for the term "eQTL" (Expression Quantitative Trait Loci) to find packages related to that topic.

Packages found under Software:

Show	All ▾	entries	Search table:	eQTL
Package	Maintainer	Title		
gQTLBase	VJ Carey	gQTLBase: infrastructure for eQTL, mQTL and similar studies		
gQTLstats	VJ Carey	gQTLstats: computationally efficient analysis for eQTL and allied studies		
iBMQ	Greg Imholte	integrated Bayesian Modeling of eQTL data		
Showing 1 to 3 of 3 entries (filtered from 1,381 total entries)			◀ Previous	Next ▶

For more information, you can check out the package list [here](#).

Whether you are using the pre-made workflows or ended up creating your own, you can likely speed up processing time by running your Bioconductor/R scripts in parallel. Microsoft R Server and `RevoScaleR` make this easy.

Let's take the [Annotating Genomic Variants](#) workflow, for example.

.vcf files are often very large and sometimes difficult to process or summarize due to their size. Using the `VariantAnnotation::locateVariants` function from Bioconductor makes this process more automated. We can use this function to identify where a variant falls with respect to gene structure, e.g., exon, utr, splice site, etc. We use the gene model from the `TxDb.Hsapiens.UCSC.hg19.knownGene` package loaded earlier.

```
## Use the 'region' argument to define the region
```

```
## of interest. See ?locateVariants for details.
cds <- locateVariants(vcf, txdb, CodingVariants())
five <- locateVariants(vcf, txdb, FiveUTRVariants())
splice <- locateVariants(vcf, txdb, SpliceSiteVariants())
intron <- locateVariants(vcf, txdb, IntronVariants())

all <- locateVariants(vcf, txdb, AllVariants())
```

If we want to start summarizing the variants, we could use `sapply` to repetitively perform some operation over the entire data object. Take a look at the highlighted lines below.

```
aa <- predictCoding(vcf, txdb, Hsapiens)
idx <- sapply(split(mcols(aa)$QUERYID, mcols(aa)$GENEID, drop=TRUE), unique)
sapply(idx, length)
```

```
## Summarize variant location by gene:
sapply(names(idx),
  function(nm) {
    d <- all[mcols(all)$GENEID %in% nm, c("QUERYID", "LOCATION")]
    table(mcols(d)$LOCATION[duplicated(d) == FALSE])
  })
```

```
##           125144 162514 23729 51393 7442 84690
## spliceSite      0      2      0      0      1      0
## intron          0      0      0      0      0      0
## fiveUTR         0      2      0      1      3      5
## threeUTR        0     25      2      1      2      0
## coding          0      5      0      3      8      0
## intergenic      0      0      0      0      0      0
## promoter        1     23      0      0     15     11
```

This code easily summarizes the .vcf file in a few seconds. However, the NA06985_17.vcf.gz file is only a small (35MB) sample from human Chromosome 17. What if you were to use multiple files to assess a sample population such as the ones available from [1000 Genomes](#)? `sapply` might take a while...

We can use the `RevoScaleR` package to parallelize the summarization function in the code. By using `rxExec` to distribute the processing over multiple cores of a processor or even multiple nodes on a Hadoop cluster, we can speed up the processing time tremendously. In the sample code below, we use the `rxExec` function to split up the processing by *GeneID*.

```
## Summarize variant location by gene using rxExec from Microsoft R Server:
vcflocationsummary <- function(nm) {
  d <- all[mcols(all)$GENEID %in% nm, c("QUERYID", "LOCATION")]
  table(mcols(d)$LOCATION[duplicated(d) == FALSE])
}
rxExec(vcflocationsummary, rxElemArg(GENEID))
```

Note: this is only sample code. To fully use this workflow, visit [Bioconductor's workflow variants](#).

Try it Out!

Now that we have explored how easy it is to speed up your genomics workflows using Microsoft R Server, you can try it out for yourself. Pick a workflow that fits your needs and then use it. Once you start seeing where the processing bottlenecks are, think about using `RevoScaleR`'s parallelization functions to speed things up. Look for loops and `apply` functions as prime candidates for distributed processing.

First time using Microsoft R Server and the `RevoScaleR`? [Microsoft's documentation](#) is a great place to start. To compare the `RevoScaleR` functions, read [Explore R and RevoScaleR in 25 functions](#). If you still have questions, please [reach out to us](#) and we will be happy to help!